Weekly Activity Report — AI Answer Safety Gate

Week ending: Friday 16 Oct 2026 (report prepared 14 Oct) **Team:** Search & Answering Infrastructure **In one sentence:** we tested four ways of checking whether an AI's retrieved documents actually contain the answer before it responds, and we now have a clear recommendation for which one to deploy.

---

1. What we're building, in plain English

You already know how RAG works: when a user asks a question, the system *retrieves* a handful of relevant documents and hands them to a large language model, which writes an answer from them. That works well — **until the retrieved documents don't actually contain the answer.** The model then either makes something up or says "I don't know," and we have no way of knowing in advance which will happen.

Our project adds a **safety gate in front of the answer**. Think of it as an airport security checkpoint for every question:

- **Green light** → the documents clearly contain the answer; answer immediately and confidently.
- **Red light** → the documents are insufficient; refuse or escalate to a human.

The gate works by asking a small, very fast AI model one question per request: "*What is the probability these documents are sufficient to answer this question?*" — and only letting answers through when that probability is high enough. "Small and fast" matters: this checkpoint runs on *every* query, so if it's slow or expensive, the whole product slows down with it.

The single number that governs everything is our risk budget: we agreed the gate may let through a confidently-wrong answer on **no more than 5% of the questions it approves**. That 5% is the "false-sufficient rate" (FPR) you'll see throughout.

2. What we did this week

1. **Finished all benchmark runs — including the complete exam.** Every candidate model was tested on the full public SufficiencyBench exam (1,390 questions, each paired with a deliberately complete and an incomplete document set, so we always know the right verdict in advance). The final full run took 78 minutes and produced **13,900 judgements with zero failures**. All automated checks pass.
2. **Fixed a serious flaw in our two-step design** (see §4, "the combo"), re-ran it, and confirmed the fix works end-to-end — and still holds at full scale.
3. **Made the deployment decision** and stress-tested it on the complete dataset: this is the week we can say, with statistics behind it, which gate to ship.
4. **Widened the talent search: eight new open-source "instant judgement" models** from around the internet were discovered, vetted, and sat the same full exam under the same rules (§10). Six

completed; one couldn't even run on our hardware; one failed an entrance exam. **Result: the crown stands, but the runner-up changes.**

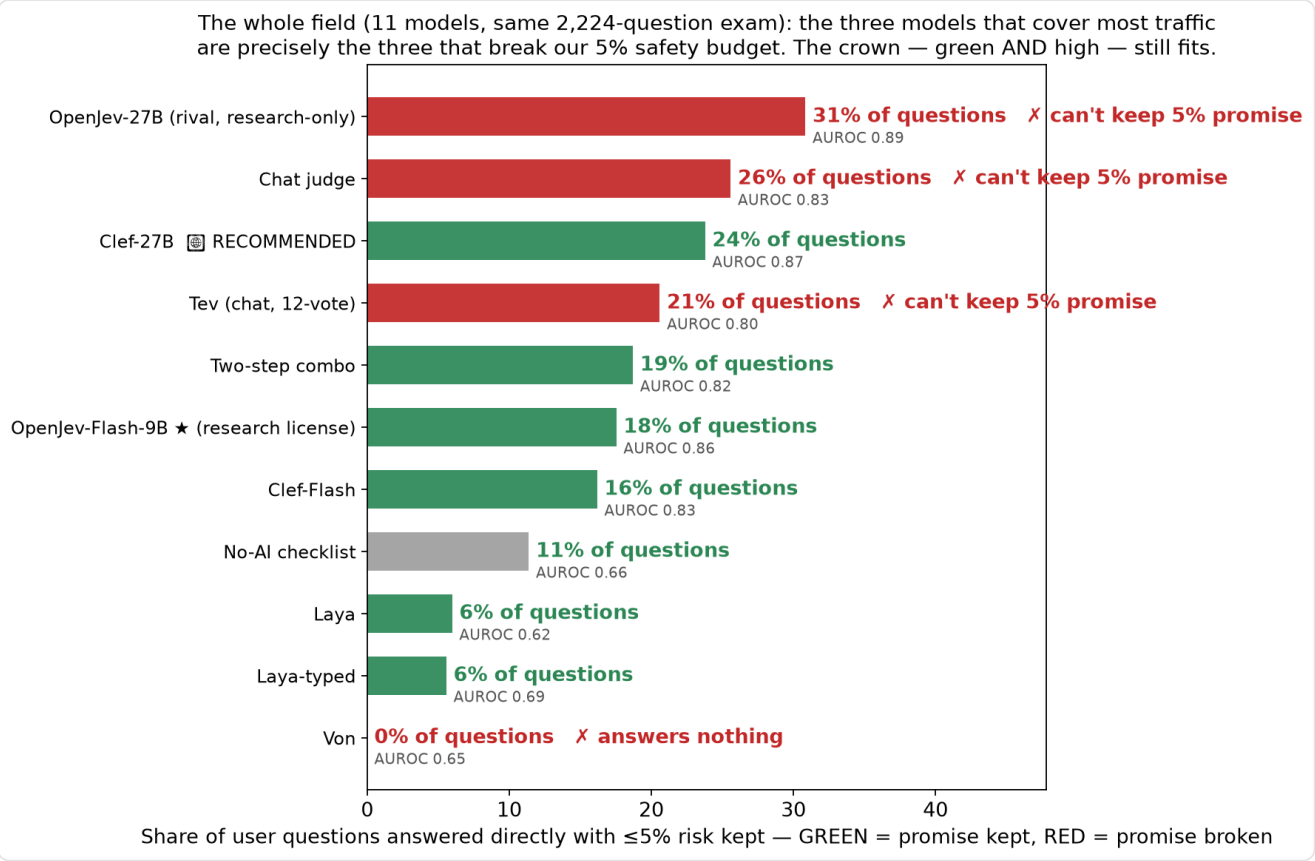
3. The five contestants

Contestant	What it is	Analogy
Clef-Flash	Small (3B) purpose-built "instant judgement" model	Sharp junior analyst — fast, usually right
Clef-27B	Same idea, 9× bigger	Senior analyst — slower, noticeably more reliable
Chat judge	General chat model asked to referee each case	Consultant — talks a good game, overconfident
Two-step combo	Junior triages everything; only the hard 22% goes to the consultant	Junior + escalation
Baselines	Simple word-overlap and size checks	Checklists, no AI at all

…and six new challengers (§10)

Challenger	What it is	Analogy
OpenJev-27B	Biggest of a rival "instant judgement" family	Brilliant but cocky senior analyst — sharpest instincts, signs off on things he shouldn't
OpenJev-Flash-9B	Mid-size from the same family	Solid senior analyst — beats our junior on the exam, works a bit slower
Tev	A chat model trained to commit to an answer without waffle	Analyst forced to answer yes/no, 12 times, no hedging
Laya & Laya-typed	Lightweight "one-glance" checkers (one pass, no thinking)	Speed-reading interns
Von	Single-glance gut-feel checker, Apache-licensed	The intern who stamps "don't know" rather than risk an error

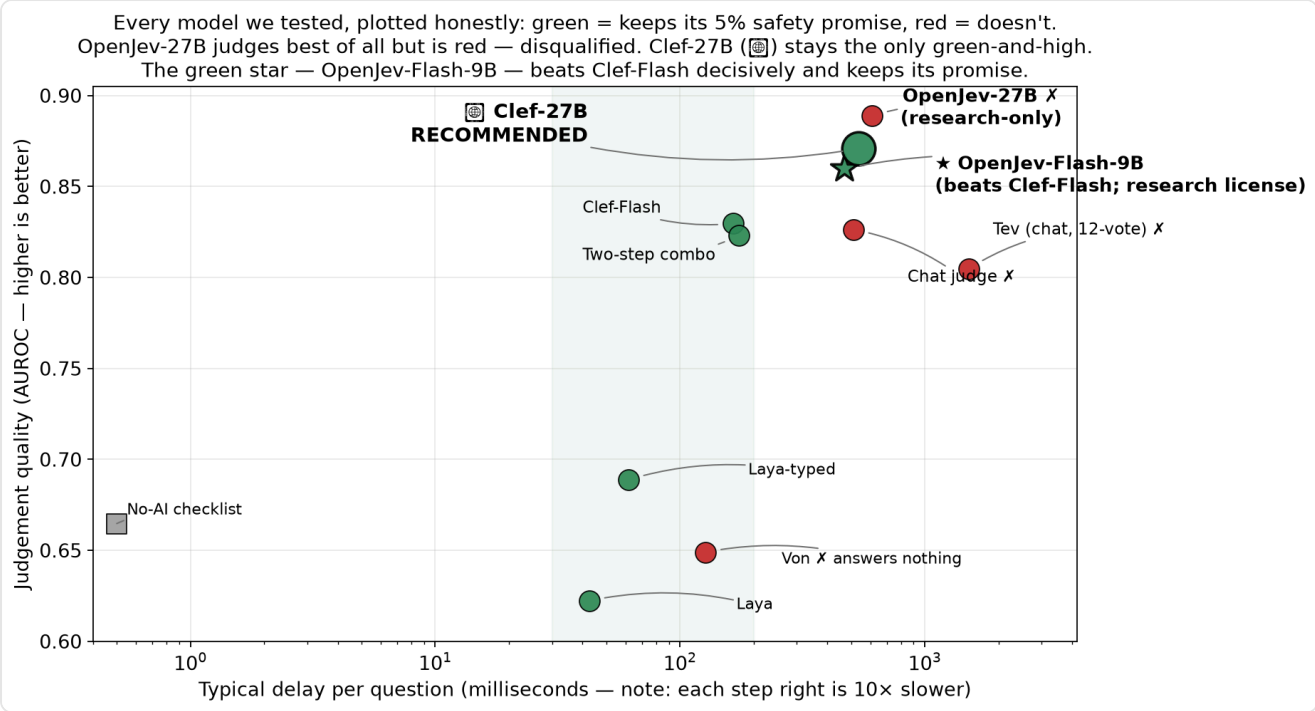
4. Results — the chart that matters



How to read this chart: each bar is one of the eleven models we've now tested; the length is the share of all user questions the gate would let the AI answer directly while keeping its error rate within our 5% budget. **Green = the model honestly keeps that promise; red = it breaks it.** Shorter bars aren't failures — the cheap interns (Laya, 6%) simply aren't smart enough to tell which questions are safe, so they approve almost nothing.

What it says (full dataset, 2,224 test questions — final numbers): - The three widest bars — OpenJev-27B (31%), chat judge (26%), Tev (21%) — are **all red**: they cover the most traffic by promising more than they can deliver. That coverage cannot be bought at the risk we accept. - **Clef-27B (green, 24%) remains the widest bar that is honest.** This week's eight-model search did not dislodge it — it confirmed that combining top judgement with top restraint is rare. - **OpenJev-Flash-9B (green, 18%) is this week's genuine catch**: a budget-compliant option that judges better than Clef-Flash — though its license (research-only, §10) keeps it out of production for now (see §10). - **The bottleneck is still timidity, not danger**: even the honest champions auto-answer only 16–24% of traffic, and roughly half of what they block was actually safe to answer. Coverage — not safety — is where the business value is locked up (§9). - **The chat judge's bar is a trap** (§6): it reaches 26% only by sounding artificially confident; on the full exam it breaks the 5% budget outright (5.5%).

5. How each candidate performed — the scoreboard



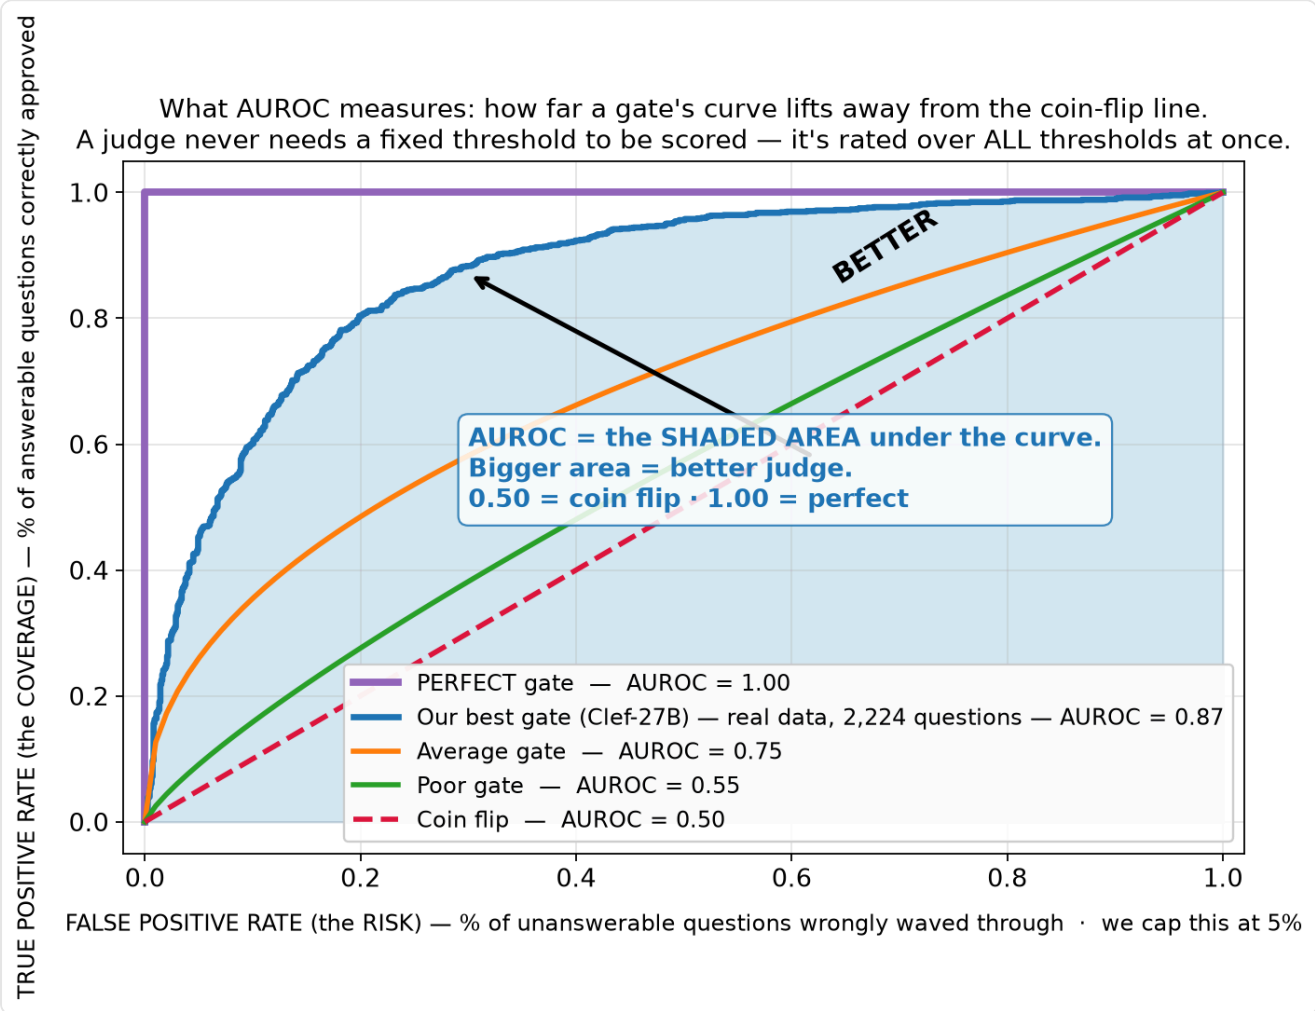
How to read this chart: every dot is one model — all eleven we've now tested. Right = higher judgement quality (AUROC, explained below); further up = lower delay per question (the axis multiplies by ten each step — Laya at the bottom-left answers 35× faster than Tev). There is no free lunch; the interesting question is who gets the best trade-off — and whether they keep their 5% safety promise (green) or break it (red).

What it says (all 11 models, full dataset): - **Clef-27B (👑) stays the pick:** the only model that is simultaneously among the best judges *and* green. Its quality lead over the whole green field is statistically decisive. - **OpenJev-27B (red, top-right) is the week's big story:** highest judgement quality we have ever measured (0.889) — but red, disqualified (§6, §10). - **OpenJev-Flash-9B (green star)** beats Clef-Flash on quality decisively at moderate speed — a real talent, but research-licensed (§10), so shelved rather than priced. - **The intern tier (bottom-left) is cheap but blunt:** Laya/Von answer almost nothing safely; Tev pays chat-model speed (far right) for mid-table green-less accuracy.

The scoreboard numbers, demystified

Term	Plain-English meaning
AUROC (0–1)	*Discrimination.* Give the judge a known-good case and a known-bad case at random — AUROC is the chance it rates the good one higher. 0.50 = coin flip; 0.88 = strong; 1.00 = perfect.
FPR (our 5% budget)	*The risk.* Of the questions the gate **approved**, how many actually had insufficient evidence — i.e. confidently wrong answers shipped to users. We cap this at 5%.
Coverage	*Throughput.* % of all traffic the gate lets through at that risk level. This is the business upside: coverage × traffic = auto-answered queries.
ECE (0–1, lower better)	*Calibration — "does its confidence match its accuracy?"* If the gate says "90% sure" on a batch of cases, it should be right ~90% of the time. High ECE means its stated confidence can't be trusted even if its gut decisions are partly right.
p50 / p95 latency	*Speed.* Typical response time (p50) and worst-case-for-95%-of-traffic (p95). We design around p95 because unhappy users remember the slow ones.
Significance (95% CI)	*Did it really win?* We re-scored the test thousands of times with reshuffled samples; if the measured gap holds in 95% of those re-runs, it's real, not luck. Clef-27B's lead is significant; the middle three are statistically tied.

What is AUROC? The 30-second version



That word — **AUROC** — is the one most likely to trip people up, so here it is with no jargon. *(The blue curve is the real, measured performance of our best gate on this week's full 2,224-question test set; the purple, orange, green and red curves are reference lines to calibrate your eye.)*

Every safety gate faces the same tension: be **strict** (few wrong answers slip through, but lots of good questions get blocked) or **relaxed** (lots of questions answered, but more wrong ones slip through). You can dial this strictness knob anywhere along a scale — a "threshold". Pick the strictest setting and you answer almost nothing; pick the loosest and you answer everything, correctly or not.

This chart shows *every possible setting of the knob at once*:

- **Right axis (the risk):** the share of unanswerable questions wrongly waved through.
- **Up axis (the payoff):** the share of answerable questions correctly approved.
- **Each point on a curve** is one strictness setting. The curve traces what happens as you slide the knob from ultra-strict (far left) to ultra-relaxed (far right).

The three reference lines:

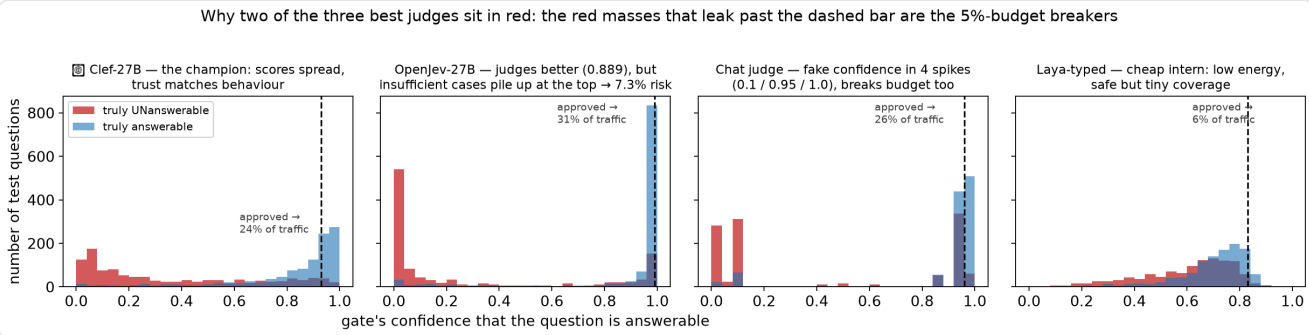
- The **purple "perfect" line** hugs the top-left corner: a flawless judge answers everything worth answering and nothing that isn't.
- The **red dashed diagonal** is the coin flip: to approve twice as many good questions, you're forced to wave through exactly twice as many bad ones. Zero skill — the payoff never outruns the risk.
- **Every real judge sits in between**, and the question is simply *how far above the coin-flip line its curve bows*.

AUROC is just the area of the shaded region under a gate's curve, as a score from 0.50 (coin flip) to 1.00 (perfect). It's the number we quote because it grades a gate's *raw judgement ability* independent of where we happen to set the threshold — the curve lifts up and left as a judge gets better, and AUROC is the single-number summary of that lift.

Two useful gut checks when you see one quoted:

- **0.50 = useless** — pure coin flip. (Our dumb control tests score exactly this, which is how we know the exam itself isn't fooling us.)
- **The gap is what matters, not the decimal.** The difference between 0.83 and 0.87 is *not* "4% better" in plain terms — in live traffic it translates directly into how many more questions the machine can be trusted to handle without a human. Clef-27B's AUROC of 0.87 vs Clef-Flash's 0.83 is a genuinely large, statistically significant capability jump — and on the full exam the statistics are now decisive, not suggestive (§5).

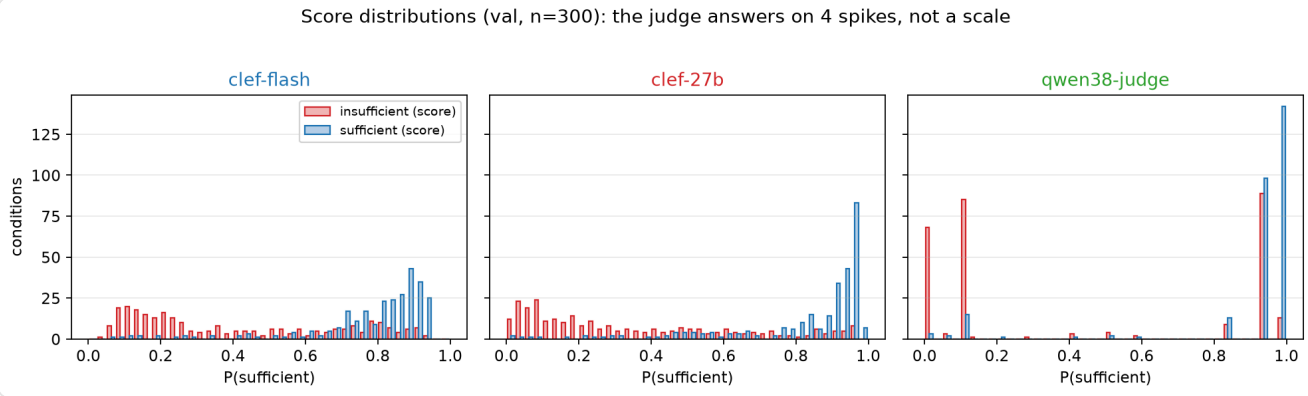
6. The week's key discovery: three models, three ways to fake confidence



How to read this chart: each panel shows how one model spread its confidence across the same 2,224 test questions — blue = cases it should approve, red = cases it should block. The dashed bar is the approval threshold. A *good* judge pushes blue right and red left with a clean gap; **any red that leaks past the dashed bar is risk we pay for.**

What it says — the new rival, the judge, and the champion: - **OpenJev-27B (second panel) fails differently from the judge — and it's the more interesting failure.** It is *not* shouty: its scores look beautifully separated (that's the 0.889). But look at the left: a big red mountain of "0–10% sure" correct rejections… and a thin but stubborn tail of red cases sitting at 95–100% that should have been blocked. Even the strictest threshold can't cut them — that tail is its 7.3% budget breach. Brilliant instincts, unlucky certainty. Our existing "translator" calibration recipe (§10 next steps) is designed to squash exactly this tail — the highest-value follow-up now on the board. - **The chat judge (third panel) shout-storm confirmed at full scale:** almost everything lands on four values (0%, 10%, 95%, 100%), and of 531 escalated cases **251 came back "95% sure" — 123 of them wrong**. Our shipped translator maps 95% → ~51% (block) and rescued the two-step combo from 2% → 19% coverage, best-calibrated on the board — still not enough to beat Clef-27B. - **Clef-27B (first panel) does the boring, correct thing:** opinions spread smoothly, no spikes, no heavy tail past the bar. Boring is exactly what a safety gate should look like. - **Laya-typed (right)** is honest but timid — its energies never reach the bar, so it approves almost nothing: safe, and nearly useless.

Historical note: the first version of this chart (below) is how we caught the judge's fake confidence last week — same phenomenon, now with the whole field to compare against.



7. What the machine costs (and it fits)

	Clef-Flash	Clef-27B	OpenJev-Flash-9B ★	OpenJev-27B (research)	Tev	Laya-typed
Typical delay per checkpoint	0.17 s	0.54 s	0.47 s	0.61 s	1.52 s (×12 votes)	0.06 s
Questions per second (serial)	~6	~2	~2	~1.6	~0.7	~16
Memory	~36 GB peak (measured)	~53 GB peak (measured)	9.3 GB weights	14.8 GB weights	9.1 GB weights	off-box sidecar (~3 GB)
Cost basis	input only, ~750 tokens/question	same	same	same	13× requests	local, no tokens

- Both Clefs fit comfortably on the Mac Studio (256 GB) **alongside** the chat models we already run (peaks measured in the full-benchmark run). The OpenJev checkpoints are physically smaller than their Clef rivals (weights listed; resident peaks during serving weren't separately captured this week — flagged for the next run), and the encoder interns run off the GPU in their own little side-cars.
- Important operational finding: **running several models at once doesn't add capacity** — throughput stays pinned and requests just queue up. We now schedule the heavy checks serially and design around ~6 questions/second per model, which changes how we size a future rollout.

8. Case files — what a real test case looks like

Numbers aside, here is what one exam question actually looks like: the same question asked about two nearly identical documents — one that contains the answer, one that doesn't. (Every example below is verbatim from this week's runs.)

Case 1 — the easy win (history)

> **Question:** "*Who was in charge of the papal army in the War of Barbastro?*"
>
> **Document A (contains the answer):** "*The legendary religious zeal of the Normans was exercised in religious wars long before the First Crusade… In 1064, during the War of Barbastro,* ***William of Montreuil led the papal army*** *and took a huge booty.*"
>
> **Document B — same document, answer sentence surgically removed:** "*…In 1018, Roger de Tosny travelled to the Iberian Peninsula to carve out a state for himself from Moorish lands, but failed.* ***By then however, there were already Norman mercenaries serving as far away as Trebizond and Georgia.***"

Gate	Doc A (answer present)	Verdict	Doc B (answer removed)	Verdict
Clef-Flash	**93%** sure	✅ answer	**11%** sure	⛔ block
Clef-27B	**97%** sure	✅ answer	**2%** sure	⛔ block
Chat judge	**100%** sure	✅ answer	**0%** sure	⛔ block

All three read the pair like humans would: *the answer is in there* vs *it isn't*. This is the gate working as designed — and the two documents are **the same length and topic** on purpose, so no model can cheat by counting words.

Case 2 — the trap, and the save (health question)

> **Question:** "*Is it possible to become physically addicted to weed?*" **Correct answer:** "*No. Marijuana is not physically addictive, but it is possible to become psychologically addicted to it.*"
>
> **Document A (sufficient):** "***No.*** *Marijuana is not physically addictive, but it is possible to become psychologically addicted to it.*"
>
> **Document B (insufficient — tricky!):** "***You absolutely could.*** *Marijuana is not physically addictive, but it is possible to become psychologically addicted to it.*"

Gate	Doc A	Doc B	What happened
Clef-Flash	90% ✅	73%	correct, twice — and flagged Doc B as *contradictory*
Chat judge	95% ✅	95%	**gave both docs the identical score** — fooled
Two-step combo	✅	⛔	judge's 95% translated to "*really only ~51% trustworthy*" → correctly blocked

This single pair is the project in miniature: one word different ("No" vs "You absolutely could"), the chat judge couldn't tell them apart, and the translation layer we built this week (§6) is what caught the trap. Clef-Flash even *diagnosed* why — it labelled Doc B "contradictory," not just "insufficient."

Case 3 — what "coverage" costs look like in practice

> **Question:** "*…determine if the value of an input is less than what variable?*" (a computer-science question about factorization)

Gate	Confidence	Threshold	Outcome
Clef-Flash	74%	89%	⛔ blocked — "not confident enough"
Clef-27B	91%	93%	⛔ blocked — *so close*
Chat judge	100%	—	saw it clearly, and the combo gate let it answer

This is where auto-answer traffic gets lost today: all three models sort of knew, no single gate crossed its bar. Whether stronger combos can safely harvest these near-misses is exactly what next week's work targets.

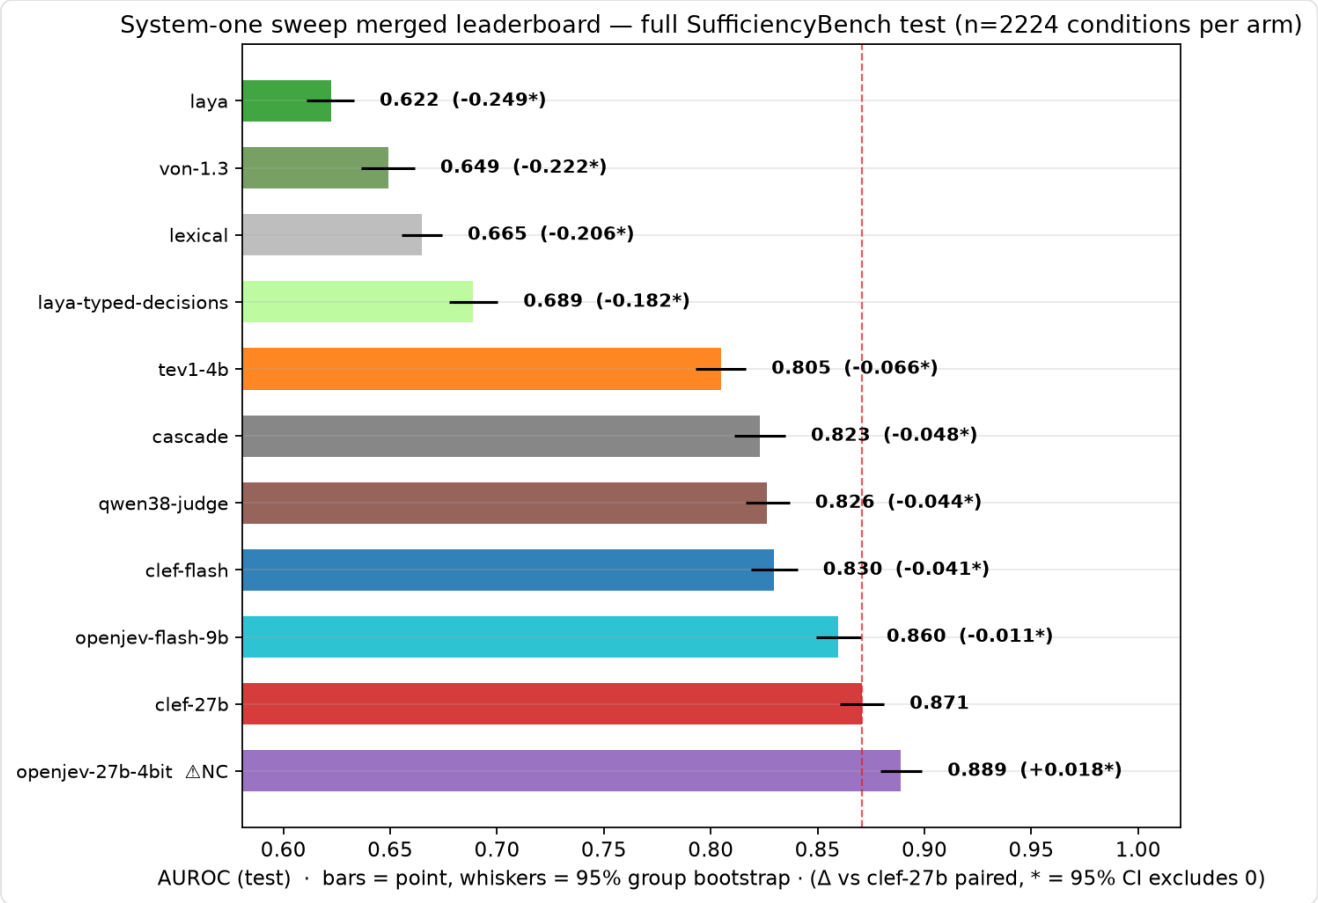
9. Recommendation & next week

Ship Clef-27B as the production safety gate. On the full exam it wins on accuracy (AUROC 0.87) and trustworthiness together; its lead over every alternative is now measured on 2,224 test questions and is statistically decisive, and the infrastructure supports it today. Keep Clef-Flash as the budget/lower-latency fallback, and park the two-step combo — correctly engineered now (it is literally the best-calibrated option on the board), but it can't yet justify its complexity. The full exam also delivered a new, clean disqualifier: the chat judge run alone **exceeds our 5% safety budget (5.5%)** — its headline coverage cannot be bought at the risk we accept. **The eight-model search (§10) did not overturn this — it strengthened it.** The one rival that judged *better* than Clef-27B broke the same safety budget the chat judge did; every other rival was further back. Clef-27B remains the only contestant that wins quality and obedience together. One concrete upgrade did fall out of the search: if/when we run the budget-compliant fast tier, **OpenJev-Flash-9B beats Clef-Flash by a statistically decisive margin** (+0.03 AUROC) at a moderate latency cost — though the license check (§10) has since come back research-only, so today the production fast tier stays Clef-Flash.

1. **Biggest opportunity:** every option today answers only ~1 in 5 questions it could safely answer. This week we found the first crack in that wall — see §11. Next we test a purpose-trained "sufficiency probe" that could lift that substantially — coverage, not safety, is where the business value is unlocked.
2. Make the chat judge *earn* its escalations (a smarter "thinking mode" for hard cases only).
3. Build a labelled contradiction test set — the gate's second job (catching documents that disagree with each other) has promising early signs but no proof yet.
4. Escalate the throughput ceiling to the platform team: today the checkpoint caps at ~6 questions/second per model; a second replica or an upstream fix (batching) unblocks rollout scale.

10. The talent search: eight new models, one verdict

This week we asked: *is there a better gate we haven't tried?* We searched Hugging Face (the public library of AI models) for every open-source model that promises the same "one-glance sufficiency judgement" our Clef models do, vetted licenses, and put every one that could run on our machine through the identical exam, identical 5% safety budget, identical statistics.



How to read this chart: every model we have ever tested, ranked by raw judgement quality (AUROC — see §5), with its honest statistical range shown as a black whisker. The red dashed line is our incumbent Clef-27B's score. A model only truly beats it if its *whole whisker* sits to the right of that line.

What the field said:

- **One rival judged better — and failed the same test that sank the chat judge.** OpenJev-27B scored a genuinely higher AUROC (0.889 — a decisive, significant win over Clef's 0.87). But its confidence runs too hot at the top: at *every* possible strictness setting it waves through 7.3% unsafe answers — over our 5% budget, no exceptions. A gate that judges brilliantly but can't be trusted to say "no" is not a gate. **The crown stands.**
- **A genuine upgrade for the budget tier — but license-blocked: OpenJev-Flash-9B.** It *does* stay inside the 5% budget, and beats Clef-Flash decisively (0.860 vs 0.830). **License check now done (2026-10-15): CC-BY-NC-4.0 — research-only, cannot ship.** The fast-tier swap is therefore off the table for production; Clef-Flash stays. (Apache-licensed re-uploads of this family exist on Hugging Face, but re-uploader license tags don't confer rights we can rely on — we won't build on those.)
- **The cheap interns are real but junior.** The two Laya models and Von run in 0.04–0.13 seconds — the cheapest AI options on the board — and Laya-typed measurably beats the no-AI checklist (a significant +0.02). But none comes close to a Clef. Von passed the safety test by refusing to approve *anything* — safe, useless.
- **A curiosity with a lesson:** Tev is a chat model trained to commit to snap answers. Quizzed 12 times per case and averaged, it scored 0.805 — respectable, but it pays full chat-model speed (1.5 seconds — 9× our junior) for mid-table accuracy.

The meta-lesson for the roadmap: **raw judgement quality is no longer scarce — reliable self-restraint is.** Two of three high-accuracy rivals self-destructed on the calibration test that Clef passes. Our next big lever (§9's sufficiency probe) should be judged on both.

11. A free lunch, technically: two disqualified models certify when they referee each other

With every model's answers already banked, we spent one afternoon (no GPUs) testing *committees* of models instead of solo artists — 84 combinations, same exam, same 5% budget, same honest protocol.

The surprise: the committee of **the chat judge + OpenJev-27B** — the two models we disqualified — **passes the safety exam (risk 4.8%) and beats the champion's judgement decisively** (quality +0.014, coverage 25% vs the champion's 24%). Why it works: a committee case is approved only if *both* members are near-certain — and the judge's habit of shouting "95%" (its fatal flaw) now acts as a **brake**: anything the judge marks 95% gets vetoed, which happens to be exactly where OpenJev's dangerous overconfidence hides. OpenJev returns the favour by vetoing the judge's own wrong "100%"s. Two flawed referees make one honest one.

Why we still don't switch: OpenJev's license is research-only (CC-BY-NC), so this committee can't ship. But the finding is worth its weight: it **proves the overconfidence tail is fixable by a second opinion**. And when we tried to teach each model to police *itself* — the solo version of this trick — it failed for every model on the board, champion included (§16 of the technical log). That's the strongest argument yet for the plan now on the roadmap: **build our own gate that carries a referee's veto in its own head** (`PLAN-PROBE.md`) — a committee in one small model we own outright.

Appendix — full technical detail for anyone who wants it: `FINDINGS.md` (15 sections of measurements and methodology; §13–§14 are this week's model search, §15 the committee result), `REPORT.md` (model-by-model summary incl. the merged 11-arm leaderboard), `PLAN-FULL-RUN.md` and `PLAN-SYSTEMONE-SWEEP.md` (how the runs were planned vs how they went), raw run data in `results/full-benchmark-v1/` (13,900 predictions) and `results/sweep-full-v1/` (16,680), the merged tables in `results/merged-leaderboard-v1/` and `results/ensemble-v1/` , all charts reproducible via `scripts/make_report_figures.py` , `scripts/make_exec_figures.py` , `scripts/make_roc_explainer.py` , `scripts/make_merged_figures.py` , `scripts/ensemble_certify.py` and `scripts/tail_calibrate.py` .*

Generated 2026-10-15 — full data: results/full-benchmark-v1 (2,780 conditions/arm, zero errors).